

A Distributed Greedy Heuristic for Computing Voronoi Tessellations With Applications Towards Peer-to-Peer Networks

Brendan Benshoof Andrew Rosen
Anu G. Bourgeois Robert W. Harrison
Department of Computer Science, Georgia State University

May 22, 2015

Outline

Background

- Motivation

- Distributed Hash Tables

DGVH

- Our Heuristic

- Peer Management

- Algorithm Analysis

Experiments

- Heuristic Accuracy

- Routing Accuracy

Conclusion

Motivation

We were creating a distributed network based off of Delaunay Triangulation and needed a fast distributed algorithm for calculating Delaunay peers. We ran into a couple of issues.

- ▶ Distributed algorithms for solving Delaunay triangulation don't really exist (every node does a global solution).
- ▶ Not any fast solutions if we start moving out of 2D Euclidean space.

This leaves approximation. However:

- ▶ Simple approximation doesn't guarantee a fully-reachable network (k -nearest neighbor, nodes in radius r).
- ▶ Other solutions [1] required prohibitively high sampling or other hidden time costs.

And nothing can handle moving nodes.

Voronoi Tessellation and Delaunay Triangulation

- ▶ Voronoi Tessellation
 - ▶ Assign each point in a space to a 'voronoi generator' to which they are closest.
- ▶ Delaunay Triangulation
 - ▶ Given a set of points connect them such that the circumcircle of each triangle does not contain any other points.

Example

content

Distributed Hash Tables

- ▶ Abstractly, a DHT is a mechanism for maintaining a large state in a decentralized network.
- ▶ In practice, the state is a large number of (*key*, *value*) records.
- ▶ A Distributed hash table assigns those records to servers and routes request for those records to those servers.
- ▶ Current incarnations of Distributed hash tables assign servers and records locations in an arbitrary metric space.
- ▶ Servers are assigned responsibility for records that are “close” to them, and to peer with “nearby” servers.
- ▶ DHTs currently use a variety metric spaces.

Applications of DHTs

- ▶ *P2P file sharing* is by far the most prominent use of DHTs. The most well-known application is BitTorrent [4].
- ▶ *Distributed Domain Name Systems* (DNS) have been built upon DHTs [5] [7]. Distributed DNSs are much more robust than DNS to orchestrated attacks, but otherwise require more overhead.
- ▶ Distributed *machine learning* [6].
- ▶ Many *botnets* are now P2P based and built using well established DHTs [8]. This is because the decentralized nature of P2P systems means there's no single vulnerable location in the botnet.

Extant Varieties of DHT

- ▶ Ring Based DHTs
 - ▶ Chord
 - ▶ Pastry
 - ▶ Tapestry
- ▶ Tree Based DHTs
 - ▶ CAN
 - ▶ Kademlia

How are DHTs and Voronoi Tessellation/Delunay Triangulation related?

A Server is responsible for records "close" to it (Voronoi Triangulation) A Server peers with other servers that bound it's Voronoi Region (Delunay Triangulation) DHTs often have peers in excess of the Delaunay triangulation to shorten lookups, however the peers that are the Delunay neighbors

- ▶ Ring Based DHTs
 - ▶ Chord: a unidirectional modulus ring metric
 - ▶ Pastry, Symphony: bidirectional modulus ring
- ▶ Tree Based DHTs
 - ▶ CAN: Euclidean distance
 - ▶ Kademlia: XOR distance

Why do we need a distributed Voronoi heuristic?

- ▶ The different topologies DHTs utilize present optimization trade-offs (lookup latency, number of lookup hops, network robustness, availability, processing overhead)
- ▶ The primary effort in implementing a new metric space in a DHT is implementing Voronoi Regions/Delunay Triangulation algorithm in that metric.
- ▶ DGVH allows many metrics to be tested without requiring the design and development effort of generating an exact Voronoi Regions/Delunay Triangulation algorithm.

So we created a Distributed Greedy Voronoi Heuristic, or DGVH for short.

Distributed Greedy Voronoi Heuristic

- ▶ Geometrically intuitive method of approximating the one-hop delaunay peers of a Node
- ▶ Is guaranteed to form a connected mesh (unlike k-nearest heuristic)
- ▶ Can be utilized in any continuous metric space

DGVH Algorithm

- 1: Given node n and its list of *candidates*.
- 2: $peers \leftarrow$ empty set that will contain n 's one-hop peers
- 3: Sort *candidates* in ascending order by each node's distance to n
- 4: Remove the first member of *candidates* and add it to *peers*
- 5: **for all** c in *candidates* **do**
- 6: m is the midpoint between n and c
- 7: **if** Any node in *peers* is closer to m than n **then**
- 8: Reject c as a peer
- 9: **else**
- 10: Remove c from *candidates*
- 11: Add c to *peers*
- 12: **end if**
- 13: **end for**

Visual Intuition

PUT PICTURES HERE!!!!

Realistic Candidate set Size

- ▶ In application, a single node will not need to calculate the triangulation for every peer in the network, rather it will only calculate it's own peers.
- ▶ A smart "join" process will prevent the Candidate Set from reaching $O(n)$
- ▶ Expected size of the candidate set is $O(\text{degree}^2)$ which compromises current peers and 2-hop peers
- ▶ The average degree of many metric spaces is $O(1)$, then therefore DGVH will practically run in $O(1)$ time in those metric spaces.

Error Mitigation

- ▶ As DGVH is a heuristic, it will likely have errors when compared to a global Delaunay triangulation.
- ▶ This error rate is difficult to find analytically and will vary between metric spaces and distributions of locations
- ▶ A simple method of mitigating this error is to keep both one and two hop peers.

Experiment 1

In our first experiment, we wanted to test the accuracy of our heuristic.

- ▶ We generated a random graph and created the Delaunay triangulation using a global solution and DGVH.
- ▶ We found DGVH had approximately one error per node.

Our second set of experiments demonstrate that this error rate is sufficiently low for distributed applications.

Results




Figure: As the size of the graph increases, we see approximately 1 error per node. We can also see that the error rate and number of nodes has a linear relationship.

Experiment 2

Our second experiment were designed to test how well nodes could form a routing topology using DGVH.

- ▶ For each trial, we create a random graph.
- ▶ During the first two cycles, nodes are given 10 random connections.
- ▶ In each cycle (including the first two):
 - ▶ Each node gossips and runs DGVH.
 - ▶ 2000 random lookups from random nodes to random locations performed.
- ▶ The rate of successful lookups approaches 1.0 as time progresses.

Results




Figure: These figures show, starting from a randomized network, DGVH forms a stable and consistent network topology. The Y axis is the percentage of successful lookups out of 2000 queries and the X axis is the number of gossip cycles.

Other Applications

One type of distributed system that can use Voronoi tessellations are *wireless ad-hoc networks*.

- ▶ Solves the coverage-boundary problem [2].
- ▶ Can be used for sleep scheduling [3]



Olivier Beaumont, A-M Kermarrec, Loris Marchal, and Etienne Rivière.

Voronet: A scalable object network based on voronoi tessellations.

In Parallel and Distributed Processing Symposium, 2007. IPDPS 2007. IEEE International, pages 1–10. IEEE, 2007.



Bogdan Carbutar, Ananth Grama, and Jan Vitek.

Distributed and dynamic voronoi overlays for coverage detection and distributed hash tables in ad-hoc networks.

In Parallel and Distributed Systems, 2004. ICPADS 2004. Proceedings. Tenth International Conference on, pages 549–556. IEEE, 2004.



Xinyu Chen, Michael R Lyu, and Ping Guo.

Voronoi-based sleeping configuration in wireless sensor networks with location error.

In Networking, Sensing and Control, 2008. ICNSC 2008. IEEE International Conference on, pages 1459–1464. IEEE, 2008.



Bram Cohen.

The bittorrent protocol specification, 2008.



Russ Cox, Athicha Muthitacharoen, and Robert T Morris.

Serving dns using a peer-to-peer lookup service.

In Peer-to-Peer Systems, pages 155–165. Springer, 2002.



Mu Li, Li Zhou, Zichao Yang, Aaron Li, Fei Xia, David G Andersen, and Alexander Smola.

Parameter server for distributed machine learning.



Vasileios Pappas, Daniel Massey, Andreas Terzis, and Lixia Zhang.

A comparative study of the dns design with dht-based alternatives.

In *INFOCOM*, volume 6, pages 1–13, 2006.



Sherif Saad, Issa Traore, Ali Ghorbani, Bassam Sayed, David Zhao, Wei Lu, John Felix, and Payman Hakimian.

Detecting p2p botnets through network behavior analysis and machine learning.

In *Privacy, Security and Trust (PST), 2011 Ninth Annual International Conference on*, pages 174–180. IEEE, 2011.